



Assignment Project Exams Help

<https://powcoder.com>

Add WeChat powcoder

Paolo Turrini

🏠 [www.dcs.warwick.ac.uk/~pturrini](http://www.dcs.warwick.ac.uk/~pturrini) ✉ [p.turrini@warwick.ac.uk](mailto:p.turrini@warwick.ac.uk)

# Assignment Project Exam Help

We have seen MDPs and how to calculate the optimal policy (VIA).

However:

- Maybe the state space is too big to do it
- Even if we do know the states we might not know how they are related.

Today we are going to see how to handle these cases, using Reinforcement Learning.

# Add WeChat powcoder

# Assignment Project Exam Help

<https://powcoder.com>

# Add WeChat powcoder



What if we don't know what game we are playing?

Play anyway and see what happens! and play as much as possible!

We can't possibly calculate everything

Assignment Project Exam Help

Game size	Board size N	$3^N$	Percent legal	Maximum legal game position ( $10^9$ )
1x1	1	3	33%	1
2x2	4	81	70%	57
3x3	9	19,683	64%	12,675
4x4	16	43,046,721	56%	24,318,165
5x5	25	$8.47 \times 10^{11}$	49%	$4.1 \times 10^{11}$
9x9	81	$4.4 \times 10^{38}$	23.4%	$1.039 \times 10^{38}$
13x13	169	$4.3 \times 10^{80}$	8.66%	$3.72497923 \times 10^{79}$
19x19	361	$1.74 \times 10^{172}$	1.196%	$2.08168199382 \times 10^{170}$
21x21	441	$2.57 \times 10^{217}$		

<https://powcoder.com>

Add WeChat powcoder

Figure: The complexity of Go

# Neural networks + tree search

Understanding the value of game positions using:

Neural Networks using pattern recognition from a database of previously played games.

Tree Search self-playing (a lot!) and estimating the value of moves,



David Silver et al.

Mastering the game of Go  
with deep neural networks and tree search

Nature, 2016.




David Silver et al.

Mastering the game of Go without human knowledge

Nature, 2017.

I'm going to only focus on how to infer value  
**without** using pre-processed information

# Assignment Project Exam Help




# Add WeChat powcoder

To evaluate intermediate game positions we play a huge number of games from then on.



# Assignment Project Exam Help




<https://powcoder.com>

Stochastic actions (possibly different at each state) four directions

Add WeChat powcoder



# Assignment Project Exam Help




<https://powcoder.com>

Stochastic actions (possibly different at each state) - four directions

Add WeChat powcoder

# Assignment Project Exam Help




<https://powcoder.com>

Stochastic actions (possibly different at each state) four directions

Add WeChat powcoder

# Assignment Project Exam Help


<https://powcoder.com>



Stochastic actions (possibly different at each state) four directions

Add WeChat powcoder

# Assignment Project Exam Help




<https://powcoder.com>

Stochastic actions (possibly different at each state) four directions

Add WeChat powcoder

# Assignment Project Exam Help




<https://powcoder.com>

Stochastic actions (possibly different at each state) four directions

Add WeChat powcoder

# Assignment Project Exam Help




<https://powcoder.com>

Stochastic actions (possibly different at each state) four directions

Add WeChat powcoder

# Assignment Project Exam Help




<https://powcoder.com>

Stochastic actions (possibly different at each state) four directions

Add WeChat powcoder

# Assignment Project Exam Help




<https://powcoder.com>

Stochastic actions (possibly different at each state) four directions

Add WeChat powcoder



# Assignment Project Exam Help



<https://powcoder.com>

Stochastic actions (possibly different at each state) four directions

Add WeChat powcoder

# Assignment Project Exam Help

This is what is known (by the agent) about the environment

- Partially observable (we know where we are, not where we will end up being)
- Markovian (past doesn't matter)
- Stochastic actions (we are not in full control of our choices)
- Discounted rewards (we might be more or less patient)

<https://powcoder.com>  
Add WeChat powcoder

# Assignment Project Exam Help

Passive reinforcement learning:

- I have a policy
- I don't know the probabilities
- I don't know the values of states
- I don't know the value of actions

<https://powcoder.com>


Active reinforcement learning:

- I don't even have a policy

Add WeChat powcoder

## Assignment Project Exam Help

<https://powcoder.com>




## Add WeChat powcoder

- I don't know the values nor the rewards
- I don't know the probabilities
- I'm gonna play anyway

**The plan:** I execute a series of trials until the end states, just like Monte-Carlo Tree Search!

# Assignment Project Exam Help

<https://powcoder.com>




**Remember:** the expected utility is the expected sum of discounted rewards under the policy

**Assume:**  $\gamma = 1$ , just to make things simple

## Assignment Project Exam Help

<https://powcoder.com>



Suppose I get these trials

## Add WeChat powcoder

$(1,1) - 0.04 \rightsquigarrow (2,1) - 0.04 \rightsquigarrow (3,1) - 0.04 \rightsquigarrow (2,1) - 0.04 \rightsquigarrow (3,1) - 0.04 \rightsquigarrow (3,2) - 0.04 \rightsquigarrow (3,3) - 0.04 \rightsquigarrow (3,4) + 1$

$(1,1) - 0.04 \rightsquigarrow (2,1) - 0.04 \rightsquigarrow (3,1) - 0.04 \rightsquigarrow (3,2) - 0.04 \rightsquigarrow (3,3) - 0.04 \rightsquigarrow (2,3) - 0.04 \rightsquigarrow (3,3) - 0.04 \rightsquigarrow (3,4) + 1$

$(1,1) - 0.04 \rightsquigarrow (1,2) - 0.04 \rightsquigarrow (1,3) - 0.04 \rightsquigarrow (2,3) - 0.04 \rightsquigarrow (2,4) - 1$

**Idea:** Frequency is the key!

Each trial provides a sample of the expected rewards for each state visited.

# Assignment Project Exam Help

When a transition occurs from state  $s$  to state  $s'$  we apply the following update:

$$v^\pi(s) = v^\pi(s) + \alpha(r(s) + \gamma v^\pi(s') - v^\pi(s))$$

where  $\alpha \in [0, 1]$  is a confidence parameter: how much we value the new information.


$\alpha$  can be the inverse of the number of times we visited a state: the more we visited, the less we want to learn.

**Notice:** rare transitions? well, they are rare.

<https://powcoder.com>  
Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>



Add WeChat powcoder



# Assignment Project Exam Help



<https://powcoder.com>

3	0	0	0	0
2	0		0	-1
1	0	0	0	0
	1	2	3	4

Initialise the values, for

- $\gamma = 1$
- deterministic agent
- $\alpha = \frac{1}{n+1}$  where  $n$  is the number of times we visited a state
- $r = 0$  everywhere but the terminal states

## Assignment Project Exam Help





Suppose we can walk (*Up, Up, Right, Right, Right*)

Apply the update to states, as you walk along:

$$v^{\pi}(s) = v^{\pi}(s) + \alpha(r(s) + \gamma v^{\pi}(s') - v^{\pi}(s))$$

## Assignment Project Exam Help





Suppose we can walk (*Up, Up, Right, Right, Right*)

Apply the update to states, as you walk along:

$$v^{\pi}(s) = v^{\pi}(s) + \alpha(r(s) + \gamma v^{\pi}(s') - v^{\pi}(s))$$

Assignment Project Exam Help





<https://powcoder.com>

I keep walking the same way...

Add WeChat powcoder

$$v^{\pi}(s) = v^{\pi}(s) + \alpha(r(s) + \gamma v^{\pi}(s') - v^{\pi}(s))$$

Assignment Project Exam Help



Again ( $Up, Up, Right, Right, Right$ )....

Add WeChat powcoder

$$v^{\pi}(s) = v^{\pi}(s) + \alpha(r(s) + \gamma v^{\pi}(s') - v^{\pi}(s))$$

# Assignment Project Exam Help

- We have a policy which we follow;
- We backpropagate the value with a Bellman-like adjustment;
- We can use a learning rate, depending on our confidence.

<https://powcoder.com>

Add WeChat powcoder

# Assignment Project Exam Help

Now we start without a fixed policy...

What the agent needs to learn is the values of the optimal policy

$$v(s) = r(s) + \gamma \max_a \sum_{s'} P(s'|s, a) v(s')$$

**Important:** we can't stick to our (locally optimal) habits,  
we need to try new stuff!

# Add WeChat powcoder

Exploration vs Exploitation

# Assignment Project Exam Help

$$Q(s, a)$$

the value of performing action  $a$  in state  $s$

<https://powcoder.com>

$$Q(s, a) = r(s) + \gamma \sum_{s'} P(s'|(s, a)) \max_{a'} Q(s', a')$$

Add WeChat [powcoder](https://powcoder.com)

is the value of performing action  $a$  in state  $s$



# Assignment Project Exam Help


$Q(s, a) \leftarrow Q(s, a) + \alpha(r(s) + \gamma \max_{a'} Q(s', a') - Q(s, a))$   
<https://powcoder.com>

It's a temporal difference learning, without fixed policy!

## Add WeChat powcoder

## Assignment Project Exam Help

- A  $2 \times 3$  grid world
- A pit, an exit and some walls are known in this grid world, but their locations are unknown
- Arrive at the exit: win; fall in the pit: die; hit a wall, suffer
- **Goal:** Get out of this maze (i.e. safely arrive at the exit) as quickly as possible



## RL components in this problem

- **State:** The agent's current location

- **Action:** LEFT, RIGHT, UP, DOWN

- **Environment Dynamics:**

- Collision results in no movement
- otherwise, move one square in the intended direction

- **Rewards:**

- normal move: -1
- hit a wall: -10
- die: -100
- exit: +100

- **Our Goal:** find the best route to exit

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

# Assignment Project Exam Help

- $\alpha = 0.5, \gamma = 0.9$
- All Q-values are initialised as 0

<https://powcoder.com>

Add WeChat powcoder




## Applying Q-Learning to The Maze Problem

# Assignment Project Exam Help

- $\alpha = 0.5$ ,  $\gamma = 0.9$
- All Q-values are initialised as 0
- Choose UP, and receive -100

<https://powcoder.com>

Add WeChat powcoder




# Assignment Project Exam Help

- $\alpha = 0.5, \gamma = 0.9$
- All Q-values are initialised as 0
- Choose *UP*, and receive -100
- update Q value:  
 $Q([0, 2], UP)$

$$= (1 - 0.5) \times 0 +$$

$$0.5 \times (-100 + 0.9 \times 0)$$

$$= -50$$



Add WeChat powcoder


## Applying Q-Learning to The Maze Problem

# Assignment Project Exam Help

- $\alpha = 0.5$ ,  $\gamma = 0.9$
- Choose *LEFT*, and receive -1

<https://powcoder.com>

Add WeChat powcoder



# Assignment Project Exam Help


- $\alpha = 0.5$ ,  $\gamma = 0.9$
- Choose *LEFT*, and receive -1
- update Q-value:

<https://powcoder.com>

$$= (1 - 0.5) \times 0 +$$

$$0.5 \times (-1 + 0.9 \times 0)$$

Add WeChat powcoder






## Applying Q-Learning to The Maze Problem

# Assignment Project Exam Help

- $\alpha = 0.5, \gamma = 0.9$
- Choose *UP*, and receive -10

<https://powcoder.com>

Add WeChat powcoder



# Assignment Project Exam Help

- $\alpha = 0.5, \gamma = 0.9$
- Choose  $UP$ , and receive -10

- update Q-value:

$$Q([0, 1] | UP)$$


$$= (1 - 0.5) \times 0 +$$

$$0.5 \times (-10 + 0.9 \times 0)$$

$$=$$

<https://powcoder.com>

Add WeChat powcoder




## Applying Q-Learning to The Maze Problem

# Assignment Project Exam Help

- $\alpha = 0.5, \gamma = 0.9$
- Choose *LEFT*, and receive -1

<https://powcoder.com>

Add WeChat powcoder



# Assignment Project Exam Help


- $\alpha = 0.5, \gamma = 0.9$
- Choose *LEFT*, and receive -1
- update Q-value:

<https://powcoder.com>

$$= (1 - 0.5) \times 0 +$$

$$0.5 \times (-1 + 0.9 \times 0)$$

Add WeChat powcoder




## Applying Q-Learning to The Maze Problem

# Assignment Project Exam Help

- $\alpha = 0.5, \gamma = 0.9$
- Choose *UP*, and receive -1

<https://powcoder.com>

Add WeChat powcoder



# Assignment Project Exam Help


- $\alpha = 0.5, \gamma = 0.9$
- Choose  $UP$ , and receive -1
- update Q-value:

<https://powcoder.com>

$$= (1 - 0.5) \times 0 +$$

$$0.5 \times (-1 + 0.9 \times 0)$$

Add WeChat powcoder




## Applying Q-Learning to The Maze Problem

# Assignment Project Exam Help

- $\alpha = 0.5, \gamma = 0.9$
- Choose *RIGHT*, and receive 100

<https://powcoder.com>

Add WeChat powcoder




# Assignment Project Exam Help

- $\alpha = 0.5, \gamma = 0.9$
- Choose *RIGHT*, and receive 100

- update Q value:  
 $Q([0,1], \text{RIGHT})$

$$\begin{aligned} &= (1 - 0.5) \times 0 + \\ &\quad 0.5 \times (100 + 0.9 \times 0) \\ &= 50 \end{aligned}$$



Add WeChat powcoder



## Applying Q-Learning to The Maze Problem

# Assignment Project Exam Help


- $\alpha = 0.5$ ,  $\gamma = 0.9$
- The next time agent visits  $[0,0]$  and performs  $UP$ :

$$Q([0,0], UP)$$

$$= (1 - 0.5) \times (-0.5) +$$

$$0.5 \times (1 - 1 + 0.9 \times 50)$$

$$= 21.75$$



Add WeChat powcoder

# Assignment Project Exam Help


- $\alpha = 0.5$ ,  $\gamma = 0.9$
- The next time agent visits  $[0,0]$  and performs  $UP$ :

$$Q([0,0], UP)$$

$$= (1 - 0.5) \times (-0.5) +$$

$$0.5 \times (1 - 1 + 0.9 \times 50)$$

$$= 21.75$$



Add WeChat powcoder

# Assignment Project Exam Help

- Quick learning speed.
- Model-free, no need to explicitly compute probabilities, or record trajectory.
- Guarantee to converge.

<https://powcoder.com>

Add WeChat powcoder

# Assignment Project Exam Help

- $\alpha$ :
  - Learning step
  - balance between existing experiences (weight:  $1 - \alpha$ ) and new observations (weight:  $\alpha$ )
- $\gamma$ : <https://powcoder.com>
  - Future discount
  - balance between current reward (weight:  $1$ ) and next  $N$  step's reward (weight:  $\gamma^N$ )
- $\epsilon$ : Add WeChat powcoder
  - indicating how 'bold' the agent is
  - balance between **exploitation** (take greedy action,  $1 - \epsilon$  chance) and **exploration** (take random action,  $\epsilon$  chance)

What we have seen so far

# Assignment Project Exam Help

- Decision making in sequential environments typical of AI practice
- Optimisation techniques (VLA)
- ... under incomplete information (Active/Passive Learning by Belmann updates)

**What next?** Population dynamics and multi-agent reinforcement learning.

## Add WeChat powcoder