



Assignment Project Exams Help

<https://powcoder.com>

Add WeChat powcoder

Paolo Turrini

🏠 www.dcs.warwick.ac.uk/~pturrini ✉ p.turrini@warwick.ac.uk

Assignment Project Exam Help

Markov Decision Processes

<https://powcoder.com>

Policies are strategies

Add WeChat powcoder

Assignment Project Exam Help

We now go back to a "typical" AI framework: Markov Decision Processes

- Plans and policies
- Optimal policies

These are "one player" games with perfect information.
Except they are not played on trees.

This (and more) in RL's chapters 17-18



Stuart Russell and Peter Norvig

Artificial Intelligence: a modern approach


2014 (3rd edition)

<https://powcoder.com>

Add WeChat powcoder


Assignment Project Exam Help

<https://powcoder.com>



- Start at the starting square
- Move to adjacent squares
- Collision results in no movement
- The game ends when we reach either goal state $+1$ or -1

Assignment Project Exam Help



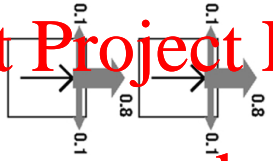
<https://powcoder.com>

The agent chooses between $\{Up, Down, Left, Right\}$ and goes:

- to the intended direction with probability: e.g., 0.8
- to the left of the intended direction with probability: e.g., 0.1
- to the right of the intended direction with probability: e.g., 0.1

Let's start walking

Assignment Project Exam Help



<https://powcoder.com>

Walking is a repetition of throws:



- The probability that I walk right the first time: 0.8
- The probability that I walk right the second time: 0.8
- The probability that I walk right both times... is a product! 0.8^2

The environment is **Markovian**: the probability of reaching a state only depends on the state the agent is in and the action they perform.

It is also **fully observable**, like an extensive game (of imperfect information).

Assignment Project Exam Help

<https://powcoder.com>





$\{Up, Down, Left, Right\}$ denote the intended directions.

A **plan** is a finite sequence of **intended** moves, **from** the start.

So $[Up, Down, Up, Right]$ is going to be the plan that, from the starting square, selects the intended moves in the specified order.

Assignment Project Exam Help

<https://powcoder.com>



Goal: get to $+1$



Consider the plan $[Up, Up, Right, Right, Right]$.

- With deterministic agents, it gets us to $+1$ with probability 1.
- But what happens to our stochastic agent instead?

What's the probability that $[Up, Up, Right, Right, Right]$ gets us to $+1$?

Assignment Project Exam Help

<https://powcoder.com>





- It's not 0.8^5 ! This is the probability that we get to $+1$ when the plan works!

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>



- It's not 0.8^5 ! This is the probability that we get to $+1$ when the plan works!
- The probability the plan does not work but still reaches $+1$ is $0.1^4 \times 0.8 = 0.00008$
- The correct answer is $0.8^5 + 0.1^4 \times 0.8$
- Notice $0.8^5 + 0.1^4 \times 0.8 < \frac{1}{3}$, not great.

Assignment Project Exam Help

S^+ is set of possible sequences of states (just like the histories of an extensive game!)

A the set of available actions.

Then a policy is a function:

<https://powcoder.com>

$$\pi : S^+ \rightarrow A$$

In words a policy is a protocol that at each possible decision point prescribes an action.


This is a strategy.

Add WeChat powcoder

A policy

Assignment Project Exam Help

<https://powcoder.com>



This is a **state-based policy**. It recommends the same action at each state (so if two sequences end up with the same state, this policy is going to recommend the same action)

Now let's complicate things a little bit...

Assignment Project Exam Help

A **reward** function is a (utility) function of the form

$r: S \rightarrow \mathbb{R}$
<https://powcoder.com>



All states, not just the terminal ones, get a reward!

Obviously, if you only care about terminal states, you may want to give zero to every other state. This is a more general model.

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>





For instance, each non-terminal state:

- has 0 reward, i.e., only the terminal states matter;
- has negative reward, e.g., each move consumes -0.04 of battery;
- has positive reward, e.g., I like wasting battery

Rewards are usually small, negative and uniform at non-terminal states. But the reward function allows for more generality.

Assignment Project Exam Help

<https://powcoder.com>



Consider now the following. The reward is:



$+1$ at state $+1$, -1 at state -1 , 0.04 in all other states.

Add WeChat powcoder

What's the expected utility of $[Up, Up, Right, Right, Right]$?

Assignment Project Exam Help

<https://powcoder.com>



Consider now the following. The reward is:

$+1$ at state $+1$, -1 at state -1 , 0.04 in all other states.



Add WeChat powcoder

What's the expected utility of $[Up, Up, Right, Right, Right]$?

IT DEPENDS

Assignment Project Exam Help

<https://powcoder.com>



Consider now the following. The reward is:

$+1$ at state $+1$, -1 at state -1 , 0.04 in all other states.

Add WeChat powcoder

What's the expected utility of $[Up, Up, Right, Right, Right]$?

IT DEPENDS on how we are going to put rewards together!

Assignment Project Exam Help

Many ways of comparing states:

- summing all the rewards
- giving priority to the immediate rewards
- ...

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

There is only one general and 'reasonable' way to combine rewards over time.

Discounted utility function: $u([s_0, s_1, s_2, \dots]) = r(s_0) + \gamma r(s_1) + \gamma^2 r(s_2) + \dots$

where $\gamma \in [0, 1]$ is the discounting factor

Notice: **additive** utility function $u([s_0, s_1, s_2, \dots]) = r(s_0) + r(s_1) + r(s_2) + \dots$ is just a discounted utility function where $\gamma = 1$.

Add WeChat powcoder

Assignment Project Exam Help

γ is a measure of the agent patience. How much more they value a gain of five today than a gain of five tomorrow, the day after etc...

- Used everywhere in AI, game theory, cognitive psychology
- A lot of experimental research on it
- Variants: discounting the discounting! I care more about the difference between today and tomorrow than the difference between some distant moment and the day after that!

<https://powcoder.com>
Add WeChat powcoder

Assignment Project Exam Help

- $\gamma = 1$ today is just another day

- $\gamma = 0$ today is all that matters

<https://powcoder.com>

Basically γ is my attitude to risk towards the future!

Notice that stochastic actions introduce further gambling into the picture

Add WeChat powcoder

A problem

Here is a 3×101 world.

50	-1	-1	-1	...	-1	-1	-1	-1
<i>s</i>				...				
50	1	1	1	...	1	1	1	1

- start at *s*.
- two deterministic actions at *s*: either *Up* or *Down*
- beyond *s* you can only go *Right*.
- the numbers are the rewards you are going to get.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

A problem

Here is a 3×101 world.

50	-1	-1	-1	...	-1	-1	-1	-1
<i>s</i>				...				
50	1	1	1	...	1	1	1	1

- start at *s*.
- two deterministic actions at *s*: either *Up* or *Down*
- beyond *s* you can only go *Right*.
- the numbers are the rewards you are going to get.

Compute the expected utility of each action as a function of γ

Assignment Project Exam Help

The utility of U_p is

$$50\gamma - \sum_{t=2}^{101} \gamma^t = 50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

The utility of Up is

$$50\gamma - \sum_{t=2}^{101} \gamma^t = 50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

The utility of $Down$ is

$$-50\gamma + \sum_{t=2}^{101} \gamma^t = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

<https://powcoder.com>
Add WeChat powcoder

Assignment Project Exam Help

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

Solving numerically, we have $\gamma \approx 0.9844$.

https://powcoder.com

Add WeChat powcoder

Assignment Project Exam Help

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

Solving numerically, we have $\gamma \approx 0.9844$.

- If γ is strictly larger than this then *Down* is better than *Up*;

Add WeChat powcoder

Assignment Project Exam Help

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

Solving numerically, we have $\gamma \approx 0.9844$.

- If γ is strictly larger than this then *Down* is better than *Up*;
- If γ is strictly smaller than this then *Up* is better than *Down*;

Add WeChat powcoder

Assignment Project Exam Help

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

Solving numerically, we have $\gamma \approx 0.9844$.

- If γ is strictly larger than this then *Down* is better than *Up*;
- If γ is strictly smaller than this then *Up* is better than *Down*;
- Else, it does not matter.

Add WeChat powcoder

Assignment Project Exam Help

A **Markov Decision Process** is a sequential decision problem for a:

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

A **Markov Decision Process** is a sequential decision problem for a:

- fully observable environment

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

A **Markov Decision Process** is a sequential decision problem for a:

- fully observable environment
- with stochastic actions

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

A **Markov Decision Process** is a sequential decision problem for a:

- fully observable environment
- with stochastic actions
- with a Markovian transition model

<https://powcoder.com>

Add WeChat powcoder


Assignment Project Exam Help

A **Markov Decision Process** is a sequential decision problem for a:

- fully observable environment
- with stochastic actions
- with a Markovian transition model
- and with discounted (possibly additive) rewards

<https://powcoder.com>
Add WeChat powcoder

Assignment Project Exam Help



Definition

Let s be a state and a an action

Model $P(s'|s, a)$ = probability that a in s leads to s'

Reward function $r(s)$ (or $r(s, a)$, $r(s, a, s')$) =

$$\begin{cases} -0.04 & \text{(small penalty) for nonterminal states} \\ \pm 1 & \text{for terminal states} \end{cases}$$

Expected utility of a policy

Assignment Project Exam Help

The expected utility (or value) of policy π , from state s is:

$$v^{\pi}(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r(s_t)\right]$$


E is the expected utility of the sequences induced by:

- the policy π (the actions we are actually going to make)
- the initial state s (where we start)
- the transition model (where we can get to)

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>



Add WeChat powcoder

- In principle, we can go on forever!
- We are going to assume we need to keep going unless we hit a terminal state (**infinite horizon assumption**)

Assignment Project Exam Help

With discounting the utility of an infinite sequence is in fact **finite**.

If $\gamma < 1$ and rewards are bounded above by r , we have:

$$u[s_1, s_2, \dots] = \sum_{t=0}^{\infty} \gamma^t r(s_t) \leq \sum_{t=0}^{\infty} \gamma^t r = \frac{r}{1-\gamma}$$

Add WeChat powcoder

Assignment Project Exam Help

An **optimal** policy (from a state) is the policy with the highest expected utility, starting from that state.

<https://powcoder.com>

$$\pi_s^* = \operatorname{argmax}_{\pi} v^{\pi}(s)$$

We want to find the optimal policy.

Add WeChat powcoder

A remarkable fact

Assignment Project Exam Help

Theorem

With discounted rewards and infinite horizon

$\pi_s^* = \pi_{s'}^*$, for each $s' \in S$
<https://powcoder.com>

This means that the optimal policy does not depend on the sequences of states, but on the states only.

In other words, the optimal policy is a state-based policy.

Add WeChat powcoder

A remarkable fact

Assignment Project Exam Help

Theorem

With discounted rewards and infinite horizon

$\pi_s^* = \pi_{s'}^*$, for each $s' \in S$
<https://powcoder.com>

This means that the optimal policy does not depend on the sequences of states, but on the states only.

In other words, the optimal policy is a state-based policy.

Idea: Take π_a^* and π_b^* . If they both reach a state c , because they are both optimal, there is no reason why they should disagree (modulo indifference!). So π_c^* is identical for both (modulo indifference!). But then they behave the same at all states!

Assignment Project Exam Help

The **value of a state** is the value of the optimal policy from that state.

But then (VERY IMPORTANT): Given the values of the states, choosing the best action is just maximisation of expected utility!

maximise the expected utility of the immediate successors

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>

3	0.812	0.998	0.912	+1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388
	1	2	3	4

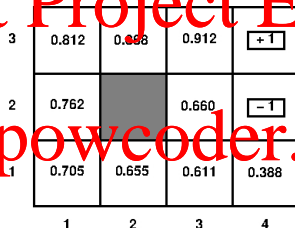
Add WeChat powcoder

Figure: The values with $\gamma = 1$ and $V(s) = -0.04$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



A 3x4 grid representing a state space. The rows are labeled 1, 2, and 3 on the left. The columns are labeled 1, 2, 3, and 4 on the bottom. The grid contains numerical values. The cell at row 2, column 2 is shaded gray. The cell at row 3, column 4 contains a box with '+1'. The cell at row 2, column 4 contains a box with '-1'.

3	0.812	0.998	0.912	<div>+1</div>
2	0.762		0.660	<div>-1</div>
1	0.705	0.655	0.611	0.388
	1	2	3	4

Figure: The optimal policy

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) v(s')$$

Maximise the expected utility of the subsequent state

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder




Figure: The optimal policy

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) v(s')$$

Maximise the expected utility of the subsequent state

Assignment Project Exam Help

The definition of values of states, i.e., the expected utility of the optimal policy from there, leads to a simple relationship among values of neighbouring states:

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

The definition of values of states, i.e., the expected utility of the optimal policy from there, leads to a simple relationship among values of neighbouring states:

<https://powcoder.com>

expected sum of rewards =
current reward $+ \gamma \times$ expected sum of rewards after taking best action

Add WeChat powcoder

Assignment Project Exam Help

Bellman equation (1957):

$$v(s) = r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a)) v(s')$$

<https://powcoder.com>

Add WeChat powcoder

The Bellman equation

Assignment Project Exam Help

Bellman equation (1957):

$$v(s) = r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a)) v(s')$$

<https://powcoder.com>

We can use it to compute the optimal policy!

Add WeChat powcoder

Assignment Project Exam Help

- 1 Start with arbitrary values
- 2 Repeat for every s simultaneously until "no change"

$$v(s) \leftarrow r(s) + \gamma \max_a \sum_{s'} v(s') P(s' | (s, a))$$

Add WeChat powcoder

Assignment Project Exam Help

- **Input** S , A , γ , r , and $P(s' | (s, a))$ for each $s, s' \in S$.
- **Input** $\epsilon > 0$, the error you want to allow
- **Output** v , the value of each state

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

① **Initialise** $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$ for all s^1 , $v := 0$, storing information to be updated

<https://powcoder.com>

Add WeChat powcoder

¹Typically this is uniform across states.

Assignment Project Exam Help

- 1 **Initialise** $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$ for all s^1 , $v := 0$, storing information to be updated
- 2 **while** $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$, for some s' , **do**

<https://powcoder.com>

Add WeChat powcoder

¹Typically this is uniform across states.

Assignment Project Exam Help

❶ Initialise $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$ for all s^1 , $v := 0$, storing information to be updated

❷ while $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$, for some s' , do

- $v(s) := r(s) + \gamma \max_{a'} \sum_{s'} P(s'|s,a') v(s')$

Add WeChat powcoder

¹Typically this is uniform across states.

Assignment Project Exam Help

- 1 **Initialise** $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$ for all s^1 , $v := 0$, storing information to be updated
- 2 **while** $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$, for some s' , **do**
 - $V(s) := r(s) + \gamma \max_{a'} \sum_{s'} P(s' | (s, a')) V(s')$
 - $\delta_s := |V'(s) - v(s)|$

<https://powcoder.com>

Add WeChat powcoder

¹Typically this is uniform across states.

Assignment Project Exam Help

❶ **Initialise** $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$ for all s^1 , $v := 0$, storing information to be updated

❷ **while** $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$, for some s' , **do**

- $V(s) := r(s) + \gamma \max_{a'} \sum_{s'} P(s' | (s, a')) V(s')$
- $\delta_s := |V'(s) - v(s)|$
- $v(s) := V'(s)$

Add WeChat powcoder

¹Typically this is uniform across states.

Assignment Project Exam Help

❶ **Initialise** $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$ for all s^1 , $v := 0$, storing information to be updated

❷ **while** $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$, for some s' , **do**

- $V(s) := r(s) + \gamma \max_{a'} \sum_{s'} P(s' | (s, a')) V(s')$
- $\delta_s := |V'(s) - v(s)|$
- $v(s) := V'(s)$

❸ **Return** v

Add WeChat powcoder

¹Typically this is uniform across states.

Assignment Project Exam Help

Theorem


<https://powcoder.com>

- *terminates*
- *returns the optimal policy (for the input values)*

Add WeChat powcoder

Assignment Project Exam Help


<https://powcoder.com>



Add WeChat powcoder

Assignment Project Exam Help


<https://powcoder.com>



Add WeChat powcoder

Assignment Project Exam Help



<https://powcoder.com>



Add WeChat powcoder

Initialise the values, for $\gamma = 1$, $r = -0.04$

Assignment Project Exam Help




Add WeChat powcoder

Simultaneously apply the Bellmann update to all states

$$v(s) = r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a)) v(s')$$

Value Iteration Algorithm



Assignment Project Exam Help

<https://powcoder.com>

3	0.812	0.868	0.912	$+1$
2	0.76		0.610	-1
1	0.705	0.655	0.611	0.388

Add WeChat powcoder

Assignment Project Exam Help


<https://powcoder.com>

$r = [-0.0480 : -0.0274]$

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>



Add WeChat powcoder

Initialise the values, for $\gamma = 1, r = -0.4$

Assignment Project Exam Help


<https://powcoder.com>

$r = [-0.4278 : -0.0850]$

Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>




Add WeChat powcoder

Initialise the values, for $\gamma = 1, r = -4$

Assignment Project Exam Help


<https://powcoder.com>



Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>




Add WeChat powcoder

Initialise the values, for $\gamma = 1, r = 0$

Assignment Project Exam Help

<https://powcoder.com>




Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Assignment Project Exam Help

- Stochastic actions can lead to unpredictable outcomes
- But we can still find optimal "strategies" exploiting what happens in case we deviate from the original plan
- If we know what game we are playing and we play long enough...

What next? Learning in MDPs

Add WeChat powcoder